

Supplementary Material

A Training-Free Quantum-Walk Scanner for Cryptic Allosteric Sites
Global Quantum + AI Challenge 2026 · Cleveland Clinic Enterprise Challenge · Phase I

1. Delivered-model evaluation

The delivered ranking is the coherent directed current on the chiral Laplacian generator, scored by the distance-deconfounded $z(j)$. All numbers below come from exact classical simulation on the 13-protein set (10 calibration + 3 blind validation); ground truth is the 4 Å heavy-atom rule. Entry 1IE4 is undefined because all its labelled residues fall inside the active-site-plus-first-neighbour exclusion, leaving zero candidate positives — an honest structural guard, not a failure.

Table 1 reports the challenge-literal test (Guidelines §4.1): one-sided Mann–Whitney U of $z(\text{distal positives})$ against (T1) random surface background and (T2) non-functional fpocket residues, as rank-biserial effect size r_{rb} and Holm-corrected p across targets. Table 2 reports blind ranking quality as distance-matched AUROC of the delivered current against the trivial –distance geometric baseline (the confound level). Results are shown in full, not filtered.

Table 1: Challenge-literal significance of the delivered (distance-corrected) $z(j)$. Bold = Holm $p < 0.05$. $r_{\text{rb}} \in [-1, 1]$: +1 = positives rank above the negative set, 0 = no difference. 2RD5 has a single distal positive (unstable).

Target	N	n_+	T1 vs background		T2 vs pockets		hit	
			r_{rb}	p_{Holm}	r_{rb}	p_{Holm}	@5	@10
Calibration (10 proteins)								
1IE4	120	0	—	—	—	—	—	—
4OFG	142	23	0.36	0.057	0.51	0.002	0	1
4B1F	255	8	0.45	0.137	0.50	0.085	1	1
5J5S	258	18	0.42	0.029	0.34	0.087	1	1
2RD5	282	1	1.00	0.102	1.00	0.087	1	1
3MWB	305	10	−0.01	1.000	−0.29	1.000	0	0
4KZT	432	14	−0.26	1.000	−0.09	1.000	0	0
1M8P	573	18	−0.39	1.000	−0.37	1.000	0	0
1YV3	703	10	−0.21	1.000	−0.30	1.000	0	0
2Y0P	847	18	−0.02	1.000	−0.01	1.000	0	0
Blind validation (3 targets)								
KRAS G12C	169	10	0.14	1.000	0.40	0.150	0	1
BCR-ABL1	451	17	0.60	<0.001	0.60	<0.001	1	1
Cardiac myosin	954	9	0.13	1.000	0.06	1.000	0	0

Table 2: Blind ranking quality across the 13-protein benchmark: distance-matched AUROC of the *delivered* coherent directed current (V5, shaded) against the single-time directed flux (V3, J) and the classical comparators — average mixing $\hat{M}$, communicability e^A , degree, and the –distance geometric baseline. **Bold** marks the delivered model above chance (> 0.5); “—” = no distal positive after filtering. Training targets are ranked by length N ; the three blind validation targets follow the rule. Classical readouts collapse to $\approx$ chance once the geometric confound is controlled.

Protein	Split	N	J (V3)	current (V5)	$\hat{M}$	e^A	degree	—dist
1IE4	train	120	—	—	—	—	—	—
4OFG	train	142	0.50	0.76	0.50	0.50	0.54	0.50
4B1F	train	255	0.35	0.70	0.44	0.32	0.23	0.74
5J5S	train	258	0.46	0.68	0.61	0.53	0.48	0.84
2RD5	train	282	0.65	1.00	0.57	0.13	0.21	0.86
3MWB	train	305	0.72	0.43	0.71	0.33	0.51	0.14
4KZT	train	432	0.54	0.43	0.34	0.43	0.66	0.63
1M8P	train	573	0.82	0.30	0.48	0.43	0.47	0.09
1YV3	train	703	0.40	0.37	0.47	0.69	0.37	0.90
2Y0P	train	847	0.30	0.47	0.39	0.42	0.42	0.32
KRAS G12C	val	169	0.60	0.67	0.34	0.33	0.33	0.54
BCR-ABL1	val	451	0.62	0.77	0.42	0.27	0.66	0.73
Cardiac myosin	val	954	0.56	0.54	0.53	0.45	0.37	0.64
Mean (12)			0.54	0.59	0.48	0.40	0.44	0.58
Mean (val)			0.59	0.66	0.43	0.35	0.45	0.64

Reading the result. On the flagship blind target BCR-ABL1 the delivered score separates the distal allosteric residues

from *both* negative sets with large effect ($r_{\text{th}} \approx 0.6$) and survives Holm correction across all targets ($p < 0.001$), and its distance-matched AUROC (0.77) beats the geometric baseline. Signal is present but weaker or absent on other targets, reported honestly; the directed current exists only through quantum coherence (Fig. 1), so it is a genuinely quantum discriminator rather than a re-expression of geometry.

2. Coherence budget and resource scaling

Figures generated by notebooks/coherence_feasibility.ipynb in the public repository, using the current best model on BCR-ABL1.




Figure 1: Coherence-time budget. (a) Systematic Trotter error falls with circuit depth r while decoherence infidelity rises; their sum is U-shaped with an optimal depth $r^* \propto \sqrt{T_2}$. (b) The ranking converges at a depth well inside the coherence-allowed r^* , so a hardware demonstration sits comfortably within the coherence budget.




Figure 2: Noise resilience and NISQ resource scaling. (a) Distal-site recovery vs the dephasing fraction p : robust across the coherent regime and collapsing to chance at the fully classical limit ($p=1$), evidence that the signal is coherence-borne; a zero-noise extrapolation recovers the coherent value. (b) Unary qubit count equals residue count and two-qubit gate count scales with contact edges: small targets are demonstrable on near-term devices, larger ones are delivered on the exact classical track.